



Institute _{of} Data

2019



Data Science and AI

Module 5

Supervised ML: Classification



Agenda: Module 5


- Introduction to **Classification**
- **Logistic Regression**
- **Evaluating** Classification Results
- **Neural Networks**
- **Support Vector Machines**
- **Bayesian Inference**
- **Applications**



Classification


Regression

- train a model by fitting data to a **continuous** response
- *predict **continuous** numbers*



Classification

- train a model by fitting data to a **discrete** response
- predict **class membership**





Logistic Regression

- Introduction to **classification**
- Logistic regression **algorithm**
- **Evaluating** classification results
- Measuring the **quality** of classification models
- **Dummy variables**



Classification

examples

- fraud detection
 - True / False
- customer segmentation:
 - frequent, high-value / regular, medium-value / occasional / transient
- credit risk
 - low / medium / high
- disease status
 - NYHA Class I / II / III / IV



Predicting Class Membership by Supervised Machine Learning

training data

- **features**
 - for now, assume these are continuous
- **response**
 - for now, assume this is binary (True/False)
 - Could be multiple classes in some case

goal

- predict $p(y = 1 \mid X)$
the probability of y being True given the predictor(s) X



Binary Class Prediction


- can we use linear regression?

$y \in \{\text{False}, \text{True}\}$



$$y = \beta X + \varepsilon$$

- binary response variable results in large residuals
- predictions can be outside $[0, 1]$





Binary Class Prediction – cont'd

- need to transform the response so that y becomes discrete
 - > model the **probability** of class membership!
- how about this:

$$p(y = 1 | X) = \beta_0 + \beta_1 X$$

- > still gives $y < 0, y > 1$
- > need an approximating function that ensures $y \in [0, 1]$



Binary Class Prediction – cont'd


- $\frac{p}{1-p}$ = odds ratio
(p = probability of True)

$$\log\left(\frac{p}{1-p}\right) = \text{logit (aka log odds)}$$

solve:

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 \mathbf{X}$$

logit function





Logistic Regression

- β_0, β_1 are known from regression results:

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X$$

let X be a new data point for prediction,

calculate \hat{y} :

$$\hat{y} = \beta_0 + \beta_1 X$$

then calculate p :

$$p = \frac{e^{\hat{y}}}{e^{\hat{y}} + 1} = \frac{1}{1 + e^{-\hat{y}}}$$

Conditions:

$$p > 0$$

$$p < 1$$




Logistic Regression

Appropriate when predicting a **binary categorical** outcome variable from a set of predictor variables (**features**) that may be **continuous and/or categorical**

Features should be independent with no missing data.

Logistic Regression typically requires a **relatively large sample size**. A general guideline is that you need at minimum of 10 cases with the least frequent outcome for each independent variable in your model.





Logistic Regression – SciKit Learn

from **sklearn.linear_model** import LogisticRegression

- Can deal with any number of features
- **Features must be numeric**
 - Categorical features should be converted to dummy features
- Can perform **multi-class classification**



Evaluating Classification Results

There are a number of metrics that can be used to evaluate a classification model. Many of these metrics revolve around values drawn from the Confusion Matrix.

The **Confusion Matrix** is a table that contains counts of the predictions of the model versus the actuals.

Note that if we are interested in predicting the opposite class, the entries would be reversed (the positive becomes negative and vice versa).

Actual	Positive	Negative
Prediction		
Positive	True Positive (TP)	False Positive (FP)
Negative	False Negative (FN)	True Negative (TN)



Confusion Matrix

It is useful to develop an intuition of the meaning of each row and column and key combinations of the counts in the Confusion Matrix

Actual	Positive	Negative	
Prediction			
Positive	True Positive (TP)	False Positive (FP)	Total predicted positive
Negative	False Negative (FN)	True Negative (TN)	Total predicted negative
	Total actual positive	Total actual negative	Total Sample count

Total true prediction





Evaluating Classification Results – cont'd

- **Accuracy**: the fraction of predictions that the model got right. i.e. Number of correct prediction / total prediction
 - Accuracy = Total true prediction / Total sample count
- **Precision**: how many of the prediction were correct. It is a measure of **exactness**.
 - Precision = TP / Total positive prediction (TP + FP)
- **Recall**: how many of the actual positive did the model predict. It is a measure of **completeness**.
 - Recall = TP / Total actual positive (TP + FN)

Actual	Positive	Negative	
Prediction			
Positive	True Positive (TP)	False Positive (FP)	Total predicted positive
Negative	False Negative (FN)	True Negative (TN)	Total predicted negative
	Total actual positive	Total actual negative	Total Sample count

Total true prediction



Actual	Positive	Negative	
Prediction			
Positive	True Positive (TP)	False Positive (FP)	Total predicted positive
Negative	False Negative (FN)	True Negative (TN)	Total predicted negative
	Total actual positive	Total actual negative	Total Sample count

Total true prediction

Note: icons are positioned where they are for illustration purpose only.

Accuracy = Total true prediction / Total sample count

Precision = TP / Total positive prediction (TP + FP)

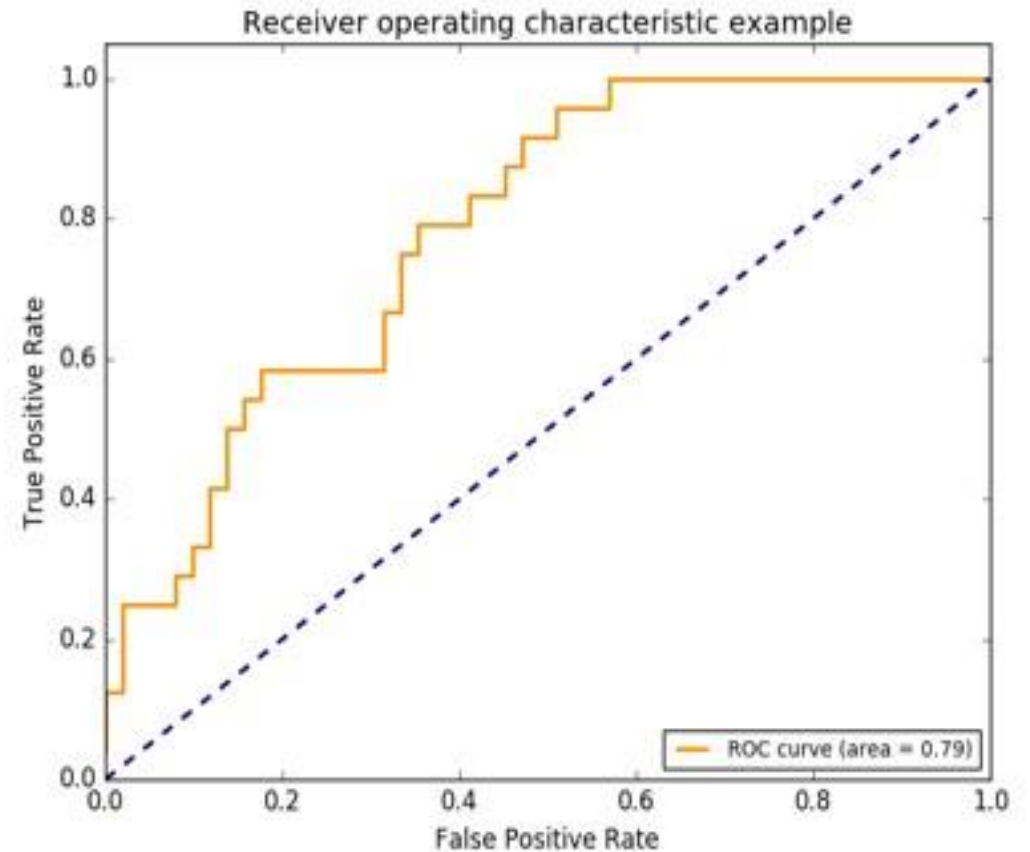
Recall = TP / Total actual positive (TP + FN)



Evaluating Classification Results – cont'd

Receiver operating characteristics (ROC) curve

- Compares True Positive Rate and False Positive Rate
- **True Positive Rate (TPR)** = Recall
- **False Positive Rate (FPR)** is FP/ Total negative count (FP+ TN)
- ROC plots **TPR** vs **FPR** by varying threshold over the entire range of threshold settings. It depicts relative **trade-offs between true positive (benefits) and false positive (costs)**
- **Area Under Curve (AUC)** is equal to the probability that the model will rank a randomly chosen positive instance higher than a randomly chosen negative one.





Evaluating Classification Results – cont'd

- There are many other metrics and many other names for the same metrics
- It is better to stick with a small number of metrics that make sense in your domain before using other metrics
- **Accuracy, precision, recall** and **AUC** are the most common metrics



Dummy Variables

How can we use categorical variables in an algorithm that requires numerical predictors?

- ***ordinal categoricals***

- can be converted to a sequence of integers, if it makes sense to do so

cold	cool	moderate	warm	hot
1	2	3	4	5



- the above implies $\text{hot} < \text{warm} < \text{moderate} < \text{cool} < \text{cold}$
... *which makes sense*



Dummy Variables

How can we use categorical variables in an algorithm that requires numerical predictors?

- *cardinal categoricals*

apples	bananas	peaches	oranges	pears
1	2	3	4	5



- > *this implies pears > oranges > peaches > bananas > apples ... does not make sense!*
- > must convert to dummy variables instead



Cardinal Dummy Variables

- full definition (number of variables = number of categories):
 - fruit_apples: 1 = apples, 0 = no apples
 - fruit_bananas: 1 = bananas, 0 = no bananas
 - fruit_peaches: 1 = peaches, 0 = no peaches
 - fruit_oranges; 1 = oranges, 0 = no oranges
 - fruit_pears; 1 = pears, 0 = no pears
- compact definition (number of variables is one less than number of categories):
 - fruit_bananas: 1 = bananas, 0 = apples
 - fruit_peaches: 1 = peaches, 0 = no peaches
 - fruit_oranges; 1 = oranges, 0 = no oranges
 - fruit_pears; 1 = pears, 0 = no pears

Python:

`pandas.get_dummies`



Lab 5.1: Logistic Regression

- Purpose:
 - To predict survival amongst Titanic passengers using the `LogisticRegression()` method of Scikit-Learn
- Materials:
 - 'Lab 5.1.ipynb'





Discussion

- Strengths & weaknesses of logistic regression




Perceptron (Neural Networks)

- **Biological** and **artificial** neurons
- **Activation** functions
- **Linear regression** with a perceptron
- **Linear classification** with a perceptron
- **Back propagation**
- **Gradient descent**
- Practical implementations



How does a nerve cell make a decision?


- Neuron receives inputs at antennae-like structures ('**dendrites**')
- Each incoming connection is **dynamically strengthened or weakened** by:
 - frequency of use ('weighting')
 - neurotransmitters
- **Weighted inputs** are summed in the cell body (transformed into a new signal)
- New signal is **propagated** along the cell's axon to be detected by other neurons





The Perceptron

- A **basic imitation** of the natural neuron
- Inputs are given **weights**
- **Weighted** signals are summed
- summed signal is transformed by an **activation function** to produce an output





Perceptron Activation Function


- z is the **weighted sum of inputs** (similar to Logistic Regression):

$$z = \sum_{j=0}^n w_j x_j$$

- A transfer function, converts z to the output of the node
- $f(z)$ is called the **activation function**


example:

$$f_{log}(z) = \frac{1}{1 + e^{-z}}$$






Perceptron Activation Functions




$$f(x) = x$$

- Linear




$$f(x) = \frac{1}{1 + e^{-x}}$$

- Logistic



$$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$$

- Hyperbolic tangent



$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$

- **Rectified linear unit (ReLU)**



Linear Regression Perceptron

- Example: given a set of fast food orders and total prices, can we predict the unit prices?

- let

x_1 = number of burgers

x_2 = number of fries

x_3 = number of sodas

y = total price


for any order

- object:

compute the weights

w_1, w_2, w_3

that minimise the residuals of y



$$y^{(i)} = w_1 x_1^{(i)} + w_2 x_2^{(i)} + w_3 x_3^{(i)}$$



Linear Classification Perceptron

- Given
 - a set of n -D **features** X
 - a set of **labels** (response var) y

- Objective:


compute the weights

$$w_1, w_2, \dots, w_n$$

that describe the **hyperplane** that separates points X by class y

example: labels $y \in \{-1, 1\}$

SEE LAB 5.2





Back Propagation

How to **compute the weights**?

- Iteratively **adjust the weights** over the entire dataset until the algorithm converges to a solution:
error (at output) < *tolerance*



How to implement back propagation?

Gradient descent

- Error at a node is a function of the weights on the inputs flowing into the node
- Use gradient of error function to step in direction of decreasing error until minimum is found
- Update weight with a Learning Rate times the error and direction of the error (which is the derivative of the cost function)
- $w = w + \text{Learning Rate} * (\text{expected} - \text{predicted}) * x$




The Perceptron in Practice

- in the *burgers/fries/sodas* example there should be an exact solution
 - (unless the unit process vary over the course of the dataset, in which case we would end up predicting the average price of each item)
- in general ...
 - the data will not permit an exact solution
 - presence of local minima may cause convergence to fail
- in practice ...
 - use many perceptrons (*nodes*), organised into *layers* to learn more complex problems
 - > neural network and Deep Learning



Lab 5.2: Classification with a Perceptron

- Purpose:
 - To evaluate a simple perceptron for predicting classes from numeric features.
- Materials:
 - 'Lab 5.2.ipynb'



Mark 1 Perceptron

By Source (WP:NFCC#4), Fair use, <https://en.wikipedia.org/w/index.php?curid=47541432>



Discussion

- Strengths & weaknesses of perceptron-based prediction




Support Vector Machines

- Concepts
- Linear SVMs
- Nonlinear SVMs
- Limitations
- Applications



Support Vector Machines


- A **linear** algebraic method for separating **n -dimensional** data into classes
 - Data points are separated by a **hyperplane** (i.e. a boundary that has dimensionality $n-1$)



- lines A, B, C are **hyperplanes** in a 2D space
- each line correctly separates the two classes
- line C is preferred, because it **maximises** the average squared distance (**margin**) between the boundary and the points




Support Vector Machines – cont'd



- line B gives a smaller margin
- line A correctly separates the classes
- line A will be chosen
 - **maximum separation** of classes is the first priority




Support Vector Machines – cont'd



- the line separates the classes except for one outlier
- the presence of the outlier does not shift the line to the right
 - **SVMs are robust against outliers**



Support Vector Machines – cont'd




- there is no line that can separate these classes
 - will an SVM fail, here?
- *Hint: we could **transform** the coordinates or create a new feature*




Support Vector Machines – cont'd

- In the original coordinates, linear separation is not possible



$z = x^2 + y^2$



- in these **coordinate**, there is a line that can separate the classes
 - > the **kernel trick**:
 - an inseparable problem can be transformed into a separable problem in a higher dimension
 - the transform function is called the **kernel**



Linear SVMs

Hard margin

- classes are separable by the margin around the solution hyperplane

Soft margin


- classes are not separable
- the best-fit hyperplane is computed by minimising the *hinge loss* function
 - incorporates a parameter that defines the trade-off between maximising the margin and minimising the number of misclassified points



Linear SVM with Hard Margin

- equation of hyperplane is
$$\vec{w} \cdot \vec{x} - b = 0$$
- compute \vec{w} , b so as to maximise width of margin

$$2/\|\vec{w}\|$$






Linear SVM with Soft Margin

- equation of hyperplane is

$$\vec{w} \cdot \vec{x} - b = 0$$

- compute \vec{w} , b so as to minimise

$$\left[\frac{1}{n} \sum_{i=1}^n \max(0, 1 - y_i(\vec{w} \cdot \vec{x}_i - b)) \right] + \lambda \|\vec{w}\|^2$$



Nb. This plot also shows that only the nearest points are used to compute the boundary. These points are called the **support vectors**.



Nonlinear SVMs

- Polynomial kernel
 - with or without constant term
- Gaussian radial basis function (RBF)
- Hyperbolic tangent



Limitations of SVMs

- Generalization error
 - margin is too small (data too close together) to accurately separate all classes
- Requires full labelling of input data
- Class membership probabilities are uncalibrated
- Does not natively handle more than 2 classes
 - multiclass problems must be reduced to a series of binary problems before applying SVMs
- Solution parameters are difficult to interpret




Applications of SVMs

- Text categorisation
- Image classification
- Handwriting interpretation
- Protein classification
- Customer segmentation



SVMs in Python

- Predictors **X**
- Response **y**
- Test data **Xtest**
- **C** = **regularization** parameter:
controls sensitivity to outliers
- **gamma** = kernel coefficient ('rbf',
'poly' and 'sigmoid' kernels):
controls influence of nearby points



```
from sklearn import svm
```

```
# Create SVM classification object:  
model = svm.svc(kernel = 'linear', c = 1, gamma = 1)
```

```
# Train model:  
model.fit(X, y)
```

```
# Evaluate quality of fit:  
model.score(X, y)
```

```
#Predict output for test data:  
predicted = model.predict(Xtest)
```



Lab 5.3: Support Vector Machines

- Purpose:
 - To apply the SVM method to linear and nonlinear classification problems.
- Materials:
 - 'Lab 5.3.ipynb'





Discussion

- Strengths & weaknesses of SVMs



Bayesian Inference

- Frequentist vs Bayesian probability
- Bayes' theorem
- Example: Disease detection
- Bayesian modelling
- Naïve Bayes Classification



Frequentist Probability

- Probability:
 - a proportion of outcomes
- In the frequentist approach, we have some **sample results** from which we calculate the **frequency** of the *positive* result
 - we estimate the future probability of a positive result based entirely on this sample



Bayesian Probability

- Probability:
 - a degree of belief
- In the Bayesian approach, we also have pre-existing information (a sample or distribution) about something that has a causal connection to the thing we are sampling
 - we use this **prior** distribution in addition to the current sample to estimate the future probability of a positive result



Bayes' Theorem – cont'd

- suppose event A leads to event B, or that B is seen to be associated with A
- let

$$p(B|A)$$

be the probability of B given A

- then

$$p(A|B) = \frac{p(A) p(B|A)}{p(B)}$$

is the probability of A given B

> *the likelihood that A must have happened given that B happened*



Bayes' Theorem – cont'd


$$p(A|B) = \frac{p(A) p(B|A)}{p(B)}$$

- $p(A)$, $p(B)$ are **marginal** probabilities
- $p(A|B)$, $p(B|A)$ are **conditional** probabilities
- $p(A)$ is the **prior probability**
- $p(B)$ is the **evidence**
- $\frac{p(B|A)}{p(B)}$ is the **support** that B provides for A



Example: Disease Detection

- prior knowledge:
 - 20% of population has the disease
 $P / (P + N) = 0.20$
 - in a random sample of 100 people,
20 will have the disease (on average)




https://arbital.com/p/bayes_frequency_diagram/



Example: Disease Detection

- sensitivity of test:
 - 90% of subjects with disease are correctly detected
$$\text{TPR} = 0.90 = \text{TP} / (\text{TP} + \text{FN})$$
$$\text{TP} = 0.90 * 20 = 18$$
- FPR (1 – specificity) of test:
 - 30% of disease-free subjects are erroneously detected
$$\text{FPR} = 0.30 = \text{FP} / (\text{FP} + \text{TN})$$
$$\text{FP} = 0.30 * 80 = 24$$




> *if a person tests positive, what is the probability that they are sick?*



Example: Disease Detection

- probability that subject with a positive test result is actually sick:

$$\begin{aligned} p(\text{sick}^+ | \text{test}^+) &= \text{TP} / (\text{TP} + \text{FP}) \\ &= 18 / (18 + 24) \\ &= 18 / 42 \\ &\approx 0.43 \end{aligned}$$



https://arbital.com/p/bayes_frequency_diagram/



Bayesian Modelling

- **Model-based** approach
 - results not dependent on arbitrary **p -value, confidence interval**
- **Naïve Bayes** assumption
 - **Predictors are independent**
 - variables come from distributions that do not interact
- prior distribution is not usually known explicitly
 - type of distribution must be assumed:
 - from domain knowledge
 - by choosing a mathematically sound distribution
 - parameters of distribution are fitted to the prior data




How to Choose the Prior Distribution?

- Informative, **empirical**:
 - Data from **related experiments** informs our prior beliefs
 - Prior beliefs will influence final predictions
- Informative, **non-empirical**:
 - Prefer certain values over others (e.g. need to **regularise coefficients**)
 - Prior beliefs will influence final predictions
- Informative, **domain-knowledge**:
 - No supporting data, but certain facts are known to be more true than others
 - Prior beliefs will influence final predictions
- **Non-informative**:
 - Prior beliefs will have little to no effect on our final assessment
 - Current data alone will determine final predictions



Beta Distribution

- a probability density function suitable for modelling the random behaviour of percentages and proportions



$$\text{PDF} = \frac{x^{\alpha-1} (1-x)^{\beta-1}}{B(\alpha, \beta)}$$

$$\alpha, \beta > 0$$

B is the beta function

https://en.wikipedia.org/wiki/Beta_distribution









Bayesian Modelling – cont'd

What can we use for the prior probability if we have no information?

> a beta distribution

$$\mu = \frac{\alpha}{\alpha + \beta}$$

$$\sigma = \sqrt{\frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}}$$






Bayesian Analysis of a Sample

Example:

Campaign to boost click rates on 'ads'



- new ad yielded 7 clicks from 10 users tested (red line)
 - beta distribution of θ peaks at 0.7 (most likely value)
- previous 100 campaigns yielded much lower click rates (light grey line)
 - prior θ peaks $\lesssim 0.17$
- we don't believe the new result, so we modify it based on our prior information (result = dark grey line)
 - posterior θ peaks $\lesssim 0.25$



Naïve Bayes Classification

- Probabilistic classification methods
- Assumptions
 - predictors are independent (hence 'naïve')
 - predictors are normally distributed
- Applications
 - text classification (spam, topic)
 - medical diagnosis
 - fraud detection
 - insurance risk category



Naïve Bayes Classification – cont'd

- Advantages
 - Algorithms **scale linearly** with number of variables
 - **Uses marginal distributions of variables**
- Disadvantages
 - **Correlated features** bias the model
 - Absolute probabilities cannot be relied upon
 - only the **probability rankings** should be used
 - Assigns zero probability if a new category appears in test data
 - training data must span all possible levels of categorical features
 - or
 - apply a smoothing technique (Sklearn uses Laplace estimation)



Naïve Bayes Classification in Python

`sklearn.naive_bayes`

GaussianNB

- assumes normally distributed features

BernoulliNB

- binary/boolean features


MultinomialNB

- discrete features (multinomial distribution)
 - e.g. word counts for text classification
- uses smoothing to account for features not present in the training set



Lab 5.4: Bayesian Inference

- Purpose:
 - To explore methods for Bayesian modelling and classification in Python
- Materials:
 - 'Lab 5.4.ipynb'
- Resources:
 - census data:
 - <https://archive.ics.uci.edu/ml/datasets/Adult>





Discussion

- more on Bayesian inference:
 - <https://www.analyticsvidhya.com/blog/2016/06/bayesian-statistics-beginners-simple-english/>
 - <https://towardsdatascience.com/bayesian-statistics-for-data-science-45397ec79c94>
- to be covered in a future module:
 - decision trees (nonlinear classification methods)



Questions?



Appendices



End of Presentation!